

# Real-Time Embedded Systems


Wan Fokkink

Vrije Universiteit Amsterdam & CWI

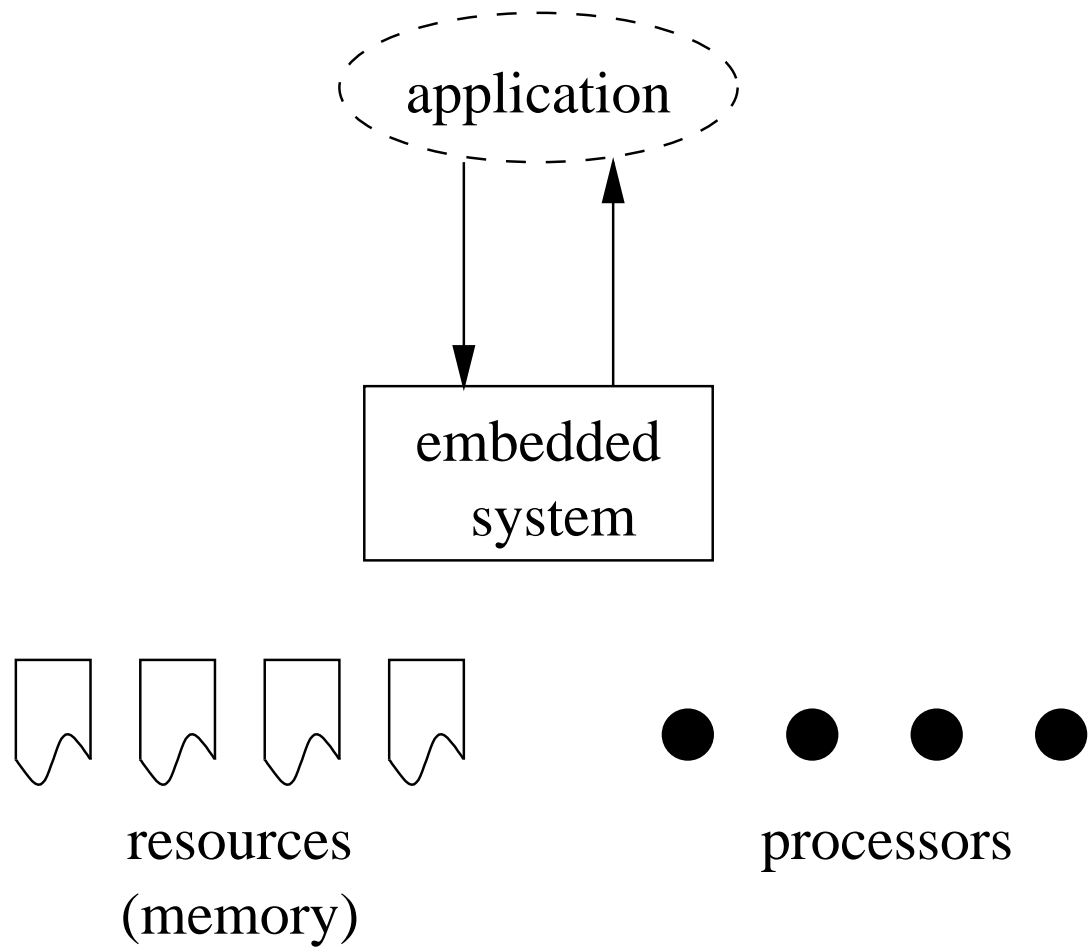
Jane W.S. Liu

*Real-Time Systems*

Prentice Hall, 2000



## General Picture



Resources are allocated to processors.

## Jobs

A **job** is a unit of work, scheduled and executed by the system.

Parameters of jobs are:

- functional behavior
- time constraints
- resource requirements

Jobs are divided over **processors**, and they are competing for

A **scheduler** decides in which order jobs are performed on a processor and which resources they can claim.

## Terminology

**release time:** when a job becomes available for execution

**execution time:** amount of processor time needed to perform (assuming it executes alone and all resources are available)

**response time:** length of time from arrival until completion

**(absolute) deadline:** when a job is required to be completed

**relative deadline:** maximum allowed response time

**hard deadline:** late completion not allowed

**soft deadline:** late completion allowed

**jitter:** imprecise release and/or execution time

A **preemptive** job can be suspended at any time of its execution

## Out of scope:

- use of distant resources
- communication between jobs
- migration of jobs
- overrun management
- penalty for missing a soft deadline
- performance
- different processor and resource types

## Types of Tasks

A **task** is a set of related jobs.

A processor distinguishes three types of tasks:

- **periodic**: known input before the start of the system, and deadlines.

Execution and interarrival times are fixed.

- **aperiodic**: executed in response to some external event, deadlines.

Execution and interarrival times are according to some *p* *distribution*.

- **sporadic**: executed in response to some external event, deadlines.

Execution times are according to some *probability distribution*, interarrival times are *random*.


## Periodic Tasks

A periodic task is defined by:

- **release time**  $r$  (of the first periodic job)
- **period**  $p$  (regular time interval, at the start of which a p is released)
- **execution time**  $e$

For simplicity we assume that the relative deadline of each p is equal to its period.

Example:  $T_1 = (1, 2, 1)$  and  $T_2 = (0, 3, 1)$ .




The conflict at time 3 is resolved by some scheduler.

The **hyperperiod** is 6.

## Job Queues at a Processor

We focus on individual aperiodic and sporadic *jobs*.



- *Sporadic jobs* are only accepted when they can be completed before their deadline.
- *Aperiodic jobs* are always accepted, and performed such that periodic and accepted sporadic jobs do not miss their deadlines.



## Average Response Time

The queueing discipline of aperiodic jobs tries to minimize e.g. **tardiness** (completion time minus deadline) or the number of soft deadlines.

The **average response time** of aperiodic jobs can be analyzed

- simulation and measurement
- Queueing Theory
- Integer Linear Programming

In the last two cases, values for e.g. average execution time of aperiodic jobs must in general still be estimated using simulation or measurement.

## Scheduler

The **scheduler** of a processor schedules and allocates resources (according to some **scheduling algorithms** and **resource access protocols**).

A schedule is **valid** if:

- jobs are not scheduled before their release times
- the total amount of processor time assigned to a job equals its (maximum) execution time

A (valid) schedule is **feasible** if all hard deadlines are met.

A scheduler is **optimal** if it produces a feasible schedule whenever possible.

## Clock-Driven Scheduler

**Off-line scheduling:** the schedule for periodic tasks is computed beforehand (typically with an algorithm for an NP-complete problem).

Time is divided into regular time intervals called **frames**.

In each frame, a predetermined set of periodic tasks is executed.

Jobs may be sliced into subjobs, to accommodate frame length.

Clock-driven scheduling is conceptually simple, but cannot cope with:

- jitter
- system modifications
- nondeterminism

## Slack

Idle time in a frame can be used to execute aperiodic and sporadic jobs.

**Slack** of a frame  $[s, c]$  at time  $t$  is  $t - c$  minus the total execution of periodic and accepted sporadic jobs in the frame after  $t$ .

**Slack stealing**: execution of aperiodic jobs until the slack of current frame is zero.

**Acceptance test**: straightforward check (in absence of jitter) whether a newly arrived sporadic job can be completed before its deadline (whether there is sufficient slack).

**Resources**: are in general distributed according to some periodic cyclic schedule.


**Example:** *Periodic* jobs  $T_1 = (0, 2, 1)$  and  $T_2 = (0, 3, 1)$ .

Frame length is 6.

*Aperiodic* job  $A$ , with execution time 2, arrives at 1.

*Sporadic* job  $S$ , with execution time 1, arrives at 2 with deadline 3.

*Sporadic* job  $S'$ , with execution time 1, arrives at 6 with deadline 7.



## Priority-Driven Scheduling

**On-line scheduling:** the schedule is computed at run-time.

Scheduling decisions are taken when:

- periodic jobs are released or aperiodic/sporadic jobs arrive
- jobs are completed
- resources are required or released

Released jobs are placed in **priority queues**, e.g. ordered by:

- release time (FIFO, LIFO)
- execution time (SETF, LETF)
- period of the task (RM)
- deadline (EDF) or slack (LST)

We focus on EDF scheduling.


## RM Scheduler

**Rate Monotonic:** Shorter period gives a higher priority.

**Advantage:** Priority on the level of tasks makes RM easier to implement than EDF/LST.

Non-optimality of the RM scheduler (one processor, preemptive, no competition for resources):

Let  $T_1 = (0, 4, 2)$  and  $T_2 = (0, 6, 3)$ .



**Remark:** If for periods  $p < p'$ ,  $p$  is always a divisor of  $p'$ , the scheduler is optimal.


## EDF Scheduler

**Earliest Deadline First:** the earlier the deadline, the higher the priority.

Consider a single processor, and preemptive jobs.

**Theorem:** When jobs do not compete for resources, the EDF scheduler is optimal.


Non-optimality in case of non-preemption:






Non-optimality in case of resource competition:

Let  $J_1$  and  $J_2$  both require resource  $R$ .



Non-optimality in case of two processors (with migration):



### Drawbacks of EDF:

- dynamic priority of periodic tasks makes it difficult to analyze which deadlines are met in case of **overloads**
- late jobs can cause other jobs to miss their deadlines (good **overrun management** is needed)

## LST Scheduler

Least Slack Time first: less slack gives a higher priority.


Slack of a job at time  $t$  is the idle time of the job until its d

Theorem: When jobs do not compete for resources, the LST is optimal.

Remarks for the LST scheduler:

- Priorities of jobs change dynamically.
- Continuous scheduling decisions would lead to context s overhead in case of two jobs with the same slack.

Non-optimality of the LST scheduler in case of two processors  
(migration):




Drawback of LST: computationally expensive

## Scheduling Anomaly

Let jobs be non-preemptive. Then shorter execution times cause violation of deadlines.

Consider the EDF (or LST) scheduler:



If jobs are preemptive, and there is no competition for resources, there is no scheduling anomaly.

## Utilization


**Utilization** of a *periodic task*  $T = (r, p, e)$  is  $\frac{e}{p}$ .

**Utilization** of a *processor* is the sum of utilizations of its periodic tasks.

Assumptions: jobs preemptive, no resource competition.

**Theorem:** Utilization of a processor is  $\leq 1$  if and only if scheduling of periodic tasks is feasible.

**Example:**  $T_1 = (1, 2, 1)$  and  $T_2 = (1, 2, 1)$ .



## Assignment of Periodic Tasks to Processors

**Goal:** To fit periodic tasks on a minimal number of processors

**Remark:** **Load balancing** is not taken into account here.

**Simple approach:** Assume processors  $P_1, \dots, P_k$ .



Periodic tasks  $T_1, \dots, T_\ell$  are assigned to processors as follows.

Let  $T_1, \dots, T_{i-1}$  have been assigned.  $T_i$  is assigned to  $P_j$  if it does not fit on  $P_1, \dots, P_{j-1}$  (i.e., utilization of these processors is beyond 1) and does fit on  $P_j$ .


**Smart approach:** First sort periodic tasks by their **utilization** ( $T_1$  has largest utilization,  $T_\ell$  smallest utilization).

This improves worst-case and average complexity (in number of required processors).

Example: Utilizations  $\frac{1}{3}$   $\frac{1}{3}$   $\frac{1}{2}$   $\frac{2}{3}$ .




simple





smart

However, the smart approach is not optimal. Fitting periodic  
a minimal number of processors is NP-complete.

Example: Utilizations  $\frac{1}{8}$   $\frac{1}{8}$   $\frac{1}{8}$   $\frac{1}{8}$   $\frac{1}{8}$   $\frac{1}{8}$   $\frac{1}{6}$   $\frac{1}{6}$   $\frac{1}{4}$   $\frac{2}{3}$ .



smart



optimal



**Remark:** Communication overhead between jobs on different processors, or between a job and a remote resource, may be account.

The problem of assigning periodic tasks to processors with m communication overhead can be reformulated into Integer Linear Programming.

## Scheduling Aperiodic Jobs


**Background:** aperiodic jobs are only scheduled in idle time.

**Drawback:** needless delay of aperiodic jobs.

**Slack stealing:** periodic tasks and accepted sporadic jobs may be interrupted if there is sufficient slack.

**Example:**  $T_1 = (0, 2, \frac{1}{2})$  and  $T_2 = (1, 3, \frac{1}{2})$ .

Aperiodic jobs available in  $[0, 6]$ .



**Drawback:** difficult to compute in case of jitter.

**Polling:** gives a **period**  $p_s$ , and an **execution time**  $e_s$  for aperiodic jobs in such a period.

At the start of a new period, *the first  $e_s$  time units* can be execute aperiodic jobs.

Consider periodic tasks  $T_k = (r_k, p_k, e_k)$  for  $k = 1, \dots, n$ . The server works if

$$\sum_{k=1}^n \frac{e_k}{p_k} + \frac{e_s}{p_s} \leq 1$$

**Drawback:** aperiodic jobs released just after a polling may be needlessly.

We proceed to present two servers based on polling that try to avoid this drawback.

For the moment, we ignore sporadic jobs.

## Deferrable Server

Allows a polling server to save its execution time within a period  $p_s$  (but not after this period!) if the aperiodic queue is empty.

In case of an EDF scheduler, the deadline of a deferrable server at the end of a period  $p_s$  can be treated as a hard deadline.


Then the deferrable server works if

$$\sum_{k=1}^n \frac{e_k}{p_k} + \frac{e_s}{p_s} \left(1 + \frac{p_s - e_s}{p_i}\right) \leq 1$$

for  $i = 1, \dots, n$  (Ghazalie & Baker, 1995).

Remark:  $\sum_{k=1}^n \frac{e_k}{p_k} + \frac{e_s}{p_s} \leq 1$  is not good enough.

Example:  $T_1 = (2, 5, 3 + \epsilon)$  and  $p_s = 3$ ,  $e_s = 1$ .



$T_1$  misses its deadline at 7

Drawback: only partial use of available bandwidth.

## Total Bandwidth Server

Fix an allowed utilization rate  $\tilde{u}_s$  for the server, such that

$$\sum_{k=1}^n \frac{e_k}{p_k} + \tilde{u}_s \leq 1$$

When the aperiodic queue is non-empty, a **deadline**  $d$  is determined for the head of the queue, according to the following rules.

(Let the head of the aperiodic queue have **execution time**  $e$ .)

- When a job arrives at the *empty* aperiodic queue at time  $t$ :

$$d := \max(d, t) + \frac{e}{\tilde{u}_s}$$


- When an aperiodic job *completes* and the *tail* of the aperiodic queue is *non-empty*:

$$d := d + \frac{e}{\tilde{u}_s}$$

Initially,  $d = 0$ .

Aperiodic jobs can now be treated as periodic jobs by the EDF scheduler.

Example:  $T_1 = (0, 2, 1)$  and  $T_2 = (0, 3, 1)$ . We fix  $\tilde{u}_s = \frac{1}{6}$ .



$J$  released at 1 with  $e = 2$  gets (at 1) deadline  $1+12=13$ .

$J'$  released at 2 with  $e' = 1$  gets (at 12) deadline  $13+6=19$ .

**Drawback:** unfair in case of multiple servers; and computationally expensive compared to the deferrable server.

**Generalized processor sharing** divides available time over servers in a round-robin fashion.

## Acceptance Test for Sporadic Jobs

A sporadic job with deadline  $d$  and execution time  $e$  is accepted if utilization (of the periodic and accepted sporadic jobs) in time interval  $[t, d]$  is never more than  $1 - \frac{e}{d-t}$ .

If accepted, utilization in  $[t, d]$  is increased with  $\frac{e}{d-t}$ .

**Example:** Periodic task  $T_1 = (0, 2, 1)$ .

Sporadic job with  $r = 1$ ,  $e = 2$  and  $d = 6$  is accepted. Utilization in  $[1, 6]$  is increased to  $\frac{9}{10}$ .

Sporadic job with  $r = 2$ ,  $e = 2$  and  $d = 20$  is rejected.

Sporadic job with  $r = 3$ ,  $e = 1$  and  $d = 13$  is accepted. Utilization in  $[3, 6]$  is increased to 1, and utilization in  $[6, 13]$  to  $\frac{3}{5}$ .



The acceptance test may reject schedulable sporadic jobs.


**Example:** Periodic task  $T_1 = (0, 2, 1)$ .

A sporadic job is released at time 0 with  $e = 1$  and  $d = 1$ .

Utilization until time 1 is 1.5, but still the sporadic job could be scheduled.

**Remark:** The total bandwidth server can be integrated with the acceptance test for sporadic jobs (e.g. by making the allowed utilization rate  $\tilde{u}_s$  dynamic.)

# Remote Access Control Protocols



Resource units can be requested by jobs during their execution. Resources are allocated to jobs in a **mutually exclusive** fashion.


When a requested resource is refused, the job is preempted (

**Remark:** Resource sharing gives rise to **scheduling anomaly**.

## Dangers of Resource Sharing


(1) Deadlock can occur.

Example:  $J_1 > J_2$ .



(2) A job  $J$  can be blocked by subsequent lower-priority jobs.

Example:  $J > J_1 > \dots > J_k$ , and  $J, J_k$  require the red resource.




## Priority Inheritance

When a job  $J$  requires a resource  $R$  and becomes blocked, the job holding  $R$  inherits the priority of  $J$  until it releases  $R$ .


(1) Deadlock can still occur.

Example:  $J_1 > J_2$ .



(2) Blocking by subsequent lower-priority jobs becomes less

Example:  $J > J_1 > \dots > J_k$ , and  $J, J_k$  require red.



## Priority Ceiling

The **priority ceiling** of a *resource*  $R$  at time  $t$  is the highest priority (known) jobs that require  $R$  at some time  $\geq t$ .

The **priority ceiling** of the *system* at time  $t$  is the highest priority ceiling of resources that are in use at time  $t$ .

(It has a special bottom value  $\Omega$  when no resources are in use.)


In a **priority ceiling protocol**, from the arrival of a job, this job is not released until its priority is higher than the priority ceiling of the system.

**Assumption:** The resources required by a job are known before its arrival.

**Note:** In the pictures to follow,  $r$  denotes the *arrival* of a job.


(1) No deadlocks. Because a job can only start executing w resources it will require are free.

Example:  $J_1 > J_2$ .



(2) Blocking by subsequent lower-priority jobs becomes less

Example:  $J > J_1 > \dots > J_k$ , and  $J, J_k$  require the red resou



This example assumes that the arrival of  $J$  is known at time

Question: What would happen if  $J$  were only known at its a

## Preemption Ceiling

**Motivation:** with dynamic priorities of tasks, the overhead of computing priority ceilings is high.

The **preemption level** of jobs must be such that if  $J > J'$ , and  $J$  arrives after  $J'$ , then  $J$  has a higher preemption level than  $J'$ .

**Idea:**  $J'$  will never preempt  $J$ .

**Example:** The preemption level can coincide with *priorities*, *arrival times*.

Ideally, the preemption level can be defined for periodic *tasks* (or jobs). For instance,

*EDF*: preemption level can coincide with RM-priorities of tasks

*FIFO*: any preemption level is allowed

*LIFO*: in general, preemption level cannot be defined for tasks

The **preemption ceiling** of a resource  $R$  at time  $t$  is the highest preemption level of (known) jobs that require  $R$  at some time.

The **preemption ceiling** of the *system* at time  $t$  is the highest preemption level of resources that are in use at time  $t$ .


In a **preemption ceiling protocol**, from the arrival of a job, the job is not released until its preemption level is higher than the preemption ceiling of the system.

*(Moreover, there is a priority inheritance rule.)*




(1) No deadlocks. Because a job can only start executing w resources it will require are free.

Example:  $J_1 > J_2$ .



(2) Blocking by subsequent lower-priority jobs becomes less

Example:  $J > J_1 > \dots > J_k$ , and  $J, J_k$  require the red resou



Note that the preemption level of  $J$  is the highest of all jobs

This example assumes that the arrival of  $J$  is known at time

## Multiple Resource Units

The notions of *priority* and *preemption ceiling* assumed only one unit per resource type.

In case of multiple units of the same resource type, the definitions of priority and preemption ceiling need to be adapted:

The **priority** (or **preemption**) **ceiling** of a resource  $R$  with  $k$  units at time  $t$  is the highest priority (or preemption level) of known jobs that require  $> k$  units of  $R$  at some time  $\geq t$ .

## Multiple Processors

In a **multiprocessor** setting, jobs that require a **global critical** resource in general are given higher priority than jobs that only require local resources.

## Greedy Synchronization Protocol

Consider a *multiprocessor* environment, where each processor runs some periodic tasks. There may be dependencies between jobs.

In the *greedy synchronization protocol*, jobs are released as soon as possible.


This schedule is not always optimal!

**Example:** Two processors  $P_1$  and  $P_2$ .

$P_1$  runs  $T_1 = (0, 6, 3)$  and  $T_2 = (0, 9, 3)$ , with  $T_1 > T_2$ .

$P_2$  runs  $T_3 = (0, 9, 3)$  and  $T_4 = (6, 10, 5)$ , with  $T_3 > T_4$ .

A job of  $T_3$  can only be released if a job of  $T_2$  completed.



**Example:** Consider the same example, but now with the EDF scheduler.

$P_1$  runs  $T_1 = (0, 6, 3)$  and  $T_2 = (0, 9, 3)$ .

$P_2$  runs  $T_3 = (0, 9, 3)$  and  $T_4 = (6, 10, 5)$ .

A job of  $T_3$  can only be released if a job of  $T_2$  completed.



All deadlines are met.


But the response time of a **pair** of jobs from  $(T_2, T_3)$  can be large, and the completion time can become unpredictable.

## Release-Guard Protocol

Suppose a job of periodic task  $T_2 = (r, p, e)$  can only be released after a job of periodic task  $T_1$  is completed.

In the **release-guard protocol**, the  $k$ -th job of  $T_2$  is released at time  $t$  with:

- $t \geq r + (k-1)p$ ;
- the  $k$ -th job of  $T_1$  completed before  $t$ ; and
- either  $T_2$ 's processor is idle at  $t$ , or  $t$  is at least  $p$  time units after the release of the  $(k-1)$ -th job of  $T_2$ .




**Example:** Two processors  $P_1$  and  $P_2$ .

$P_1$  runs  $T_1 = (0, 3, 1)$  and  $T_2 = (0, 4, 1)$ .


$P_2$  runs  $T_3 = (0, 4, 1)$  and  $T_4 = (1, 3, 2)$ .

A job of  $T_3$  can only be released if a job of  $T_2$  completed.

First consider the greedy synchronization protocol, with EDF



Now consider the release-guard protocol, with  $T_1 > T_2$  and  $T_3 > T_4$



## Summary:

- clock-driven vs. priority-driven scheduling
- fixed priority of RM vs. dynamic priority of EDF/LST
- optimality of EDF/LST vs. non-optimality of RM
- divide periodic tasks over processors
- scheduling of aperiodic jobs  
(*slack stealing / polling / total bandwidth*)
- acceptance test for sporadic jobs (*based on utilization*)
- resource sharing  
(*priority inheritance / priority ceiling / preemption c*)
- dependencies between jobs  
(*greedy synchronization / release-guard*)